

Kaizen: Building a Performant Blockchain System Verified for Consensus and Integrity

Faria Kalim[†], Karl Palmskog^{*}, Jayasi Mehar[‡], Adithya Murali[†],
P. Madhusudan[†] and Indranil Gupta[†]

[†]University of Illinois at Urbana-Champaign


^{*}KTH; work done while at UT Austin and UIUC

[‡]Facebook; work done while at UIUC



THE UNIVERSITY OF
TEXAS
— AT AUSTIN —

Blockchains and Cryptocurrencies




Consensus Protocol Challenges

Distributed protocols need to handle:

- communication delays (asynchrony)
- node crashes, corruption
- message drops, duplication, forging

Protocol implementation challenges:

- conformance to protocol specification
- node-local performance
- absence of bugs compromising safety



Consensus System Formal Verification

Project	Paper	Protocol	Tool	LOC
Disel	POPL '18	2-phase commit	Coq	5k+
Verdi Raft	CPP '16	Raft	Coq	50k+
Velisarios	ESOP '18	PBFT	Coq	50k+
<u>Ironfleet</u>	SOSP '15	Paxos	<u>Dafny</u>	20k+
<u>Toychain</u>	CPP '18	proof-of-X	<u>Coq</u>	10k+

Interactive vs. Mostly-Automated Verification

Coq proof assistant

- much training required
- + explicit proofs
- + many libraries
- purely functional (extraction to OCaml/Haskell)


Dafny verification environment

- + less training required
- implicit proofs
- few libraries
- + functional & imperative (C# code generation)

Our Contributions

- novel combination of **Coq & Dafny** to build performant and verified blockchain system, Kaizen
- methodology based on continuous refinement
 - adapted & instantiated Coq model
 - translated Coq code (not proofs) to Dafny imperative code
 - refined C# code and linked to network shim
- performed evaluation measuring minting and consensus time


Methodology Overview*



*system is fully verified until Stage III

Stage I: Modeling and Verification Using Coq

- 1 encode system in higher-order functional language (Gallina)
- 2 prove specification **interactively** using powerful **tactics**
- 3 check soundness of every low-level step



Toychain Examples

```
Record Block := mkB { prevBlockHash : Hash;  
  txs : seq Transaction; proof : VProof }.
```

```
Record State := mkS { id: Address; peers: seq Address;  
  forest: map Hash Block; txpool: seq Transaction }.
```

```
Definition valid_chain_block (bc:seq Block) (b:Block):=  
  VAF (proof b) bc (txs b) && all [pred t | txValid t bc] (txs b).
```

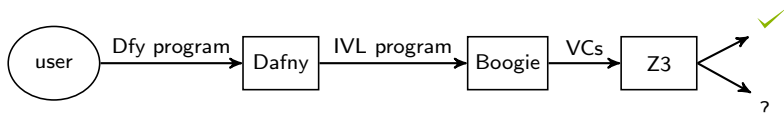
Toychain Results and Extensions

- Toychain proves that in quiescent state, all nodes know the same (canonical) chain
- we added support for coinbase transactions
- we added checking of proof-of-work validity of chains
- we changed Toychain nodes to avoid unnecessary messages

All changes are proof-preserving and now merged into Toychain.

Stage II: Refinement and Verification Using Dafny

- 1 encode programs and their contracts in imperative language
- 2 try to prove automatically that contracts are fulfilled
- 3 add more annotations if necessary



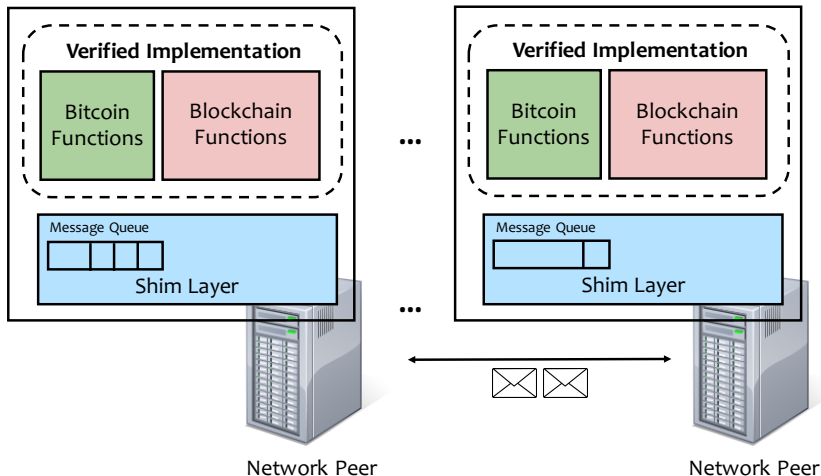
Dafny Examples

```
datatype Block = Block(prevBlockHash: Hash,  
  txs: seq<Transaction>, proof: VProof)  
  
datatype State = Node(id: Address, peers: seq<Address>,  
  forest: map<Hash,Block>, txpool: seq<Transaction>)  
  
class StateImpl {  
  var id : Address;  
  var peers :...; var forest :...; var txpool :...;  
  ghost var st: State;  
  
  predicate Valid() { ...}  
  
  method ProcMsgImpl(from: Address, msg: Message, ts: Timestamp)  
    returns (pt: seq<Packet>)  
  requires Valid();  
  ensures Valid();  
  ensures st =procMsg(old(st), from, msg, ts).0;  
  ensures pt =procMsg(old(st), from, msg, ts).1;  
  { ...}  
}
```

Stage III: Refinements in C#

- block and proof-of-work generation
- define and inject miner rewards
- store pre-computed chains
- add network shim based on UDP


Implementation Architecture




Evaluation Setup

- metrics: block minting time and consensus time
- use 30-node cluster of 2.4GHz processors w/ 64GB RAM
- baseline: performance of stock Bitcoin implementation
- workload: traces of arrival times of 50 transactions from realistic dataset


Evaluation: Block Minting Time



Evaluation: Consensus Time



Evaluation: Scalability



Components and Effort

Component	Lines of Code
Coq refinement	$\approx 1k$
Dafny refinement	$\approx 5k$
C# refinement	$\approx 1k$
C# network shim	$\approx 4k$

Development effort ≈ 6 person months across four people

Lessons Learned and Future Work

- holistic expertise necessary in Coq/Dafny/systems for Kaizen
- “easy” change can require large changes at earlier stages
- local node computation took most effort to optimize (rather than network messaging)
- future Coq proofs of Toychain Byzantine tolerance transferrable to Kaizen (see WIP by Gopinathan and Sergey, CoqPL '19)

Conclusion

- system development methodology combines interactive and mostly-automated verification, **Coq & Dafny**
- verified executable blockchain system **Kaizen**
- evaluation gives encouraging results on performance

More information:

- GitHub: <https://github.com/palmskog/kaizen>
- contact me: **Karl Palmskog** palmskog@kth.se